

Exploring Variational Autoencoders

Dami Akinniyi | July 2024

Abstract

Like many data scientists, my initial introduction to machine learning revolved around learning packages and understanding which functions to call to train a model. As a master's student applying machine learning to projects with tight deadlines, I somewhat accepted the black-box approach to understanding these models. Over the years, my priorities have shifted, and I now find myself eager to dissect the black boxes of AI to grasp the underlying mechanisms of the many models integrating into our lives.

While older than many of the generative models available today, variational autoencoders [VAEs] have a particular charm. Whether it's their rigorous and detailed formulation or the wealth of information contained in related papers, VAEs and their subsequent families are as fascinating as they are intimidating.

My first foray into VAEs involved running and modifying a hand-me-down code. After several rounds of training, it became clear that I was out of my depth and uncertain about which changes led to improvements or setbacks. After setting it aside for a while, I've finally returned to look under the hood and understand the intricacies of VAEs.

*An Introduction to Variational Autoencoders by Diederik P. Kingma and Max Welling
Autoencoding Variational Bayes by Diederik P. Kingma and Max Welling*

1 Introduction to Variational Autoencoders

Variational autoencoders are $x^2 + y^2 = z^2$

2 Understanding Variational Inference

What is this Write about concepts here. Starting a new paragraph here. Lorem ipsum. What is the understanding of this topic. I need to know what this is about and consolidate my knowledge and understanding of it. Without consistent support, it becomes rather pointless that the formatting and what is not. I don't know what to do here because this is not so great for me. Why are these paragraphs indented? I think this is a rather silly behaviour of latex.

2.1 Posterior Probability and Intractability

What does this have to do with anything

3 Code Exploration

The loss function

- minimizing evidence of the data and why mse works for fashion mnist
- in the same manner, why categorical cross entropy works for mnist digit
- Understanding expected behaviour of kl-divergence

What I have written in code and my understanding of it.